

De-anonymization and Denial-Of-Service Attacks Against the Tor Network

Indrajeet Aditya Roy

December 10, 2021

Abstract

The increasing prevalence of extensive internet surveillance has heightened the importance of secure and anonymous internet use. In response, innovative networking solutions like The Onion Router (Tor) have been developed. Tor is a low latency network that allows for anonymous and secure communication of Transmission Control Protocol (TCP) traffic. It operates on the principles of Onion Routing and decentralization. The Tor protocol uses end-to-end encryption to protect TCP streams as they pass through public networks. This encryption involves wrapping the traffic in encrypted units called Tor cells. These cells are encrypted multiple times at different points in the network, ensuring security at each transit stage. However, the decentralized nature of Tor, while vital to its function, also introduces potential risks. Its design allows any internet user to join the network, which could include malicious actors. These individuals could exploit the network by either compromising legitimate nodes or establishing harmful ones. While decentralization is key to Tor's operation, it also opens up possibilities for security vulnerabilities.

1. Introduction

As large-scale internet surveillance becomes more prevalent, concerns about the secure and anonymous use of the internet have increased. This has led to the development of innovative networking technologies, most notably The Onion Router (Tor). Tor is a low latency network designed for the anonymization of TCP-traffic communication, founded on the principles of Onion Routing and decentralization.

The Tor network ensures secure and anonymous transmission of TCP traffic through its Tor protocol. This protocol, based on Onion Routing, is an end-to-end encryption method that anonymizes and secures TCP streams over public networks. It does this by encapsulating TCP traffic in Tor cells, which are then encrypted at multiple layers at each hop in a decentralized network of routers.

While the decentralized architecture of the Tor network is one of its key strengths, it also poses significant challenges. By design, a decentralized network allows any individual on the internet to become part of the network. In terms of security, this decentralization can be a drawback for Tor, as it potentially allows malicious actors to compromise the network, either by hijacking legitimate nodes or by operating malicious nodes.

2. Network Architecture and Operation

2.1 Tor Network Components

- **Tor Cell:** Cells are a 512-byte transmission structure used in Tor, which consist of a payload and a header. Tor consists of two cell types, being control and relay.

Control cells are used to construct a transmission circuit through the network. Control cell commands are CREATE (establish a new circuit), CREATED (circuit has been successfully established) and DESTROY (delete a circuit).




Figure 1: Components of the Tor network

Relay cells are used in end-to-end transmission of TCP traffic via the transmission circuit. RELAY cell commands are BEGIN, DATA, END, SENDME, EXTEND, DROP, and RESOLVE.

- **Tor Circuit:** A multi-hop path consisting of multiple Tor nodes which facilitates anonymous low latency TCP stream from the Tor client to client destination through the Tor network.
- **Tor Relay/Node:** A volunteer operated publicly listed server, which facilitates the forwarding of traffic across the Tor network. Additionally, Tor Relays are also known as Nodes.

The Tor network consists of three relay types: Entry/Guard Node, Middle Node and Exit Node. Each Node has varying privileged access to crucial identifying information points such as IP Address of the client, client destination, and other network nodes.

Based on the principle of Onion Routing, a Node cannot view the instructions or transmitted payload of another Node. The transmitted data and instructions of a specific Tor Node can only be viewed via decryption with a shared symmetric key negotiated between the Tor Node and the client during the TLS session during path/circuit establishment.

- **Entry/Guard Node:** The entry point of a Tor client to the network. Guard Nodes are privileged nodes as they are the only point in the Tor Network where the original IP Address of the client can be observed via legitimate non-malicious processes and operation of the network.
- **Middle Node:** The primary traffic forwarding relay in the Tor Network. Majority of the Tor network consists of multiple Middle Nodes' encrypting and forwarding traffic. Middle Node's only have access to the IP addresses of the immediate predecessor node and successor node in the network.
- **Exit Node:** The primary and final exit point in a Tor circuit. Exit Nodes are privileged nodes as they are the only point in the network where the IP address of the client destination can be observed. The final transmission of the payload with the destination IP address is in plain-text but can be encrypted using TLS/SSL. As the final payload transmission is in plain-text, exit nodes are often considered the origin IP address of the traffic, consequently hiding the IP Address of the Tor client.
- **Directory Authority Node:** Directory authorities are nodes that maintain the consensus of the Tor network

by regularly updating the list of active nodes. Tor clients and Nodes are continually updated with the state of the network and the Tor Consensus by directory authority nodes. Directory Authority Node keep track of the Tor network state by periodically receiving state and operating updates from all network Nodes.

- **Tor consensus:** The Tor network uses a consensus protocol to ensure the stability and security of the network. Tor consensus is a regularly updated list of Tor nodes that are added and removed from the network. When all 9 directory authority nodes exchange, negotiate, compile and sign node information and network statuses from all the Tor directory authority nodes, a Tor consensus is formed.
- **Tor Consensus weight:** A value which determines the probability of a Tor Node being chosen by the Tor Node Selection Algorithm during the process of constructing a transmission path through the Tor network. The Tor Consensus weight of a Node is determined by the Nodes bandwidth. A large bandwidth would result in the Node being assigned a large consensus weight by the Directory Authority Node.

2.2 Tor Network Operation

1. Circuit Creation

The first step in data transmission through the Tor network involves establishing a secure circuit, initiated by the Tor client. This process involves two main components: Directory Authority nodes and the Tor Node Selection Algorithm.

The Tor Node Selection Algorithm primarily selects nodes based on bandwidth being a primary factor. Nodes with higher bandwidth have a greater probability to be chosen in the circuit creation to ensure efficient data transmission. However, the selection process also considers other factors such as:

- **Diversity in Network:** To avoid potential bottlenecks and surveillance, the algorithm selects nodes from different network families and geographical locations.
- **Entry Guards:** The Tor client typically selects a few stable and reliable nodes as its entry guards and uses them for multiple circuits to avoid exposure to malicious entry nodes.
- **Exit Policies:** Nodes have different exit policies regarding the type of traffic they allow. The algorithm selects an exit node that permits the required type of network traffic for the user's request.

Before finalizing the node selection, the Tor client verifies the nodes against the latest Tor consensus. This consensus, provided by the Directory Authorities, contains updated information about the nodes, including their status (e.g., active, down), flags (e.g., Guard, Exit, Stable), and their performance metrics. This step ensures that the nodes chosen are trustworthy and active at the time of circuit creation.

2. Circuit Encryption

The encryption process of the circuit is essentially ECDHE key exchange incrementally across the transmission circuit, negotiating TLS sessions with each Node, resulting in a shared secret key.

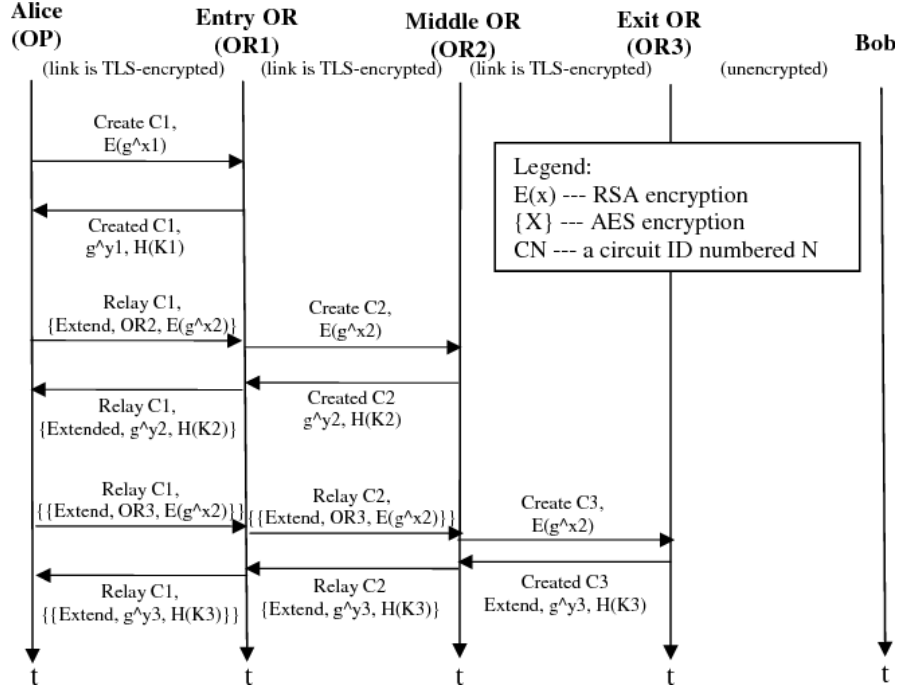


Figure 2: Circuit encryption

The initial key negotiation in the encryption process occurs between the Tor client and the Entry Node. The client (Alice) sends a **CREATE cell (C1)** to the **Entry Node (OR1)**, containing the payload $E(g^{x1})$. Here, g^{x1} represents the first part of a Diffie-Hellman (DH) handshake, and $E()$ denotes RSA encryption using the Entry Node's public onion key. In response, **OR1** sends a **CREATED cell (C1)** back to the client with the second part of the DH handshake g^{y1} and a hash of the final key $H(K1)$, after decrypting the RSA encryption with its private key. Consequently, the client and **OR1** establish a shared key, enabling **OR1** to decrypt and encrypt data during transmission.

The second key negotiation involves the Tor client and a Middle Node. The client sends a **RELAY EXTEND (C1)** cell to the **Entry Node (OR1)**. This cell includes the address of the subsequent **Middle Node (OR2)** and the payload $E(g^{x2})$. Here, g^{x2} is the initial part of a DH handshake, and $E()$ represents RSA encryption with the Middle Node's public onion key. Upon receiving the **RELAY EXTEND cell (C1)**, **OR1** extracts **OR2's** public key and forwards it to **OR2** in a **CREATE cell (C2)**. **OR2** then sends a **CREATED cell (C1)** to **OR1**, containing the second part of the DH handshake g^{y2} and a hash of the final key $H(K2)$, after decrypting the RSA encryption with its private key. Once **OR1** receives **OR2's CREATED cell**, it forwards the cell to the client in a **RELAY_EXTENDED cell (C1)**. This establishes a shared key between the **client** and **OR2**, which is used for decrypting and encrypting data during transmission.

3. Circuit Transmission

The data transmission process through a Tor circuit utilizes **RELAY DATA** cells. The client encapsulates the TCP data into the cell payload, and each **RELAY DATA** cell is encrypted with the shared symmetric key for each Node in the circuit. At each circuit hop, the Node decrypts the cell using the shared key previously negotiated between the client and the Node.

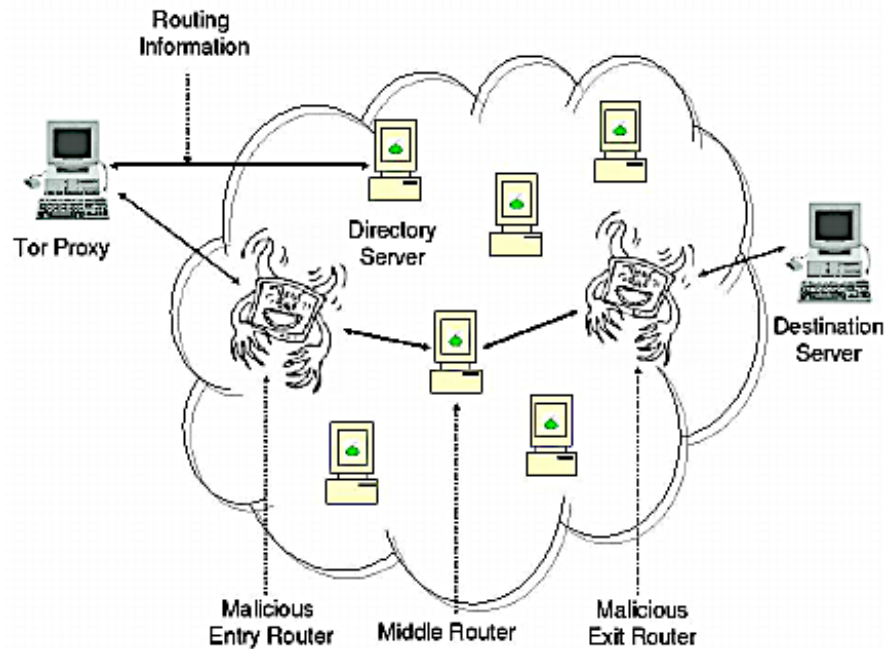


Figure 3: Sybil Attack

3. Vulnerabilities, Exploits and Mitigation Techniques

3.1 Sybil Attack

1. Description

A Sybil attack is essentially a consensus-based attack on a decentralized network, wherein an attacker is attempting to shift the network consensus in their favor by obtaining disproportionately large control in a network via the hijacking of legitimate machines or the deployment of bots or malicious machines. The attack essentially exploits the one of the primary bases of decentralization, being that any individual can become a component of the network.

2. Vulnerability

In order to initiate the encrypted TCP stream transmission, the Tor client initially constructs a transmission circuit through the Tor network via with the Tor Node Selection Algorithm which essentially optimizes the circuit for efficient transmission by choosing high bandwidth Nodes. The Tor client chooses the Nodes by referencing the Tor consensus which is periodically updated by Tor Directory Authority Nodes.

The consensus weight of a Tor Node is essentially a direct correlation to the bandwidth of the Node. A Tor Node with a high bandwidth will have a high consensus weight. A high consensus weight will increase the probability that the Node will be chosen by Tor Node Selection Algorithm, consequently allowing the Tor Node to observe more TCP traffic across the network.

3. Attack Vector

An attacker can potentially increase the consensus weight of his nodes and successfully control and observe more traffic across the tor network. The attacker can set up a network of high bandwidth, high uptime Sybil nodes or increase the bandwidth of a single Node. A single Sybil Node with a usually high amount of traffic and a large bandwidth may alert the Tor Directory Authority Nodes and allow the Tor Directory Authority Nodes to more easily remove the malicious Node from the Tor consensus. Additionally, the attacker can setup multiple Sybil Nodes that share a single IP address.[8][9]

A malicious Sybil Node being chosen as a Middle Node in a Tor circuit may allow the attacker to execute de-anonymization attacks such as Website fingerprinting and Bridge Address Harvesting.

- **Website Fingerprinting**

A website fingerprinting attack essentially enables the attacker to trace back a specific destination IP address to the source IP address, consequently de-anonymizing the client's activity. Even though the Middle Nodes in a Tor circuit cannot see the encrypted TCP stream of the entire Tor circuit, an attacker can potentially make use of packet instructions and information transmitted via the Node to form a correlation between the origin IP Address and destination IP Address.

- **Bridge Address Harvesting**

A Bridge address harvesting attack essentially enables the attacker to identify Bridge nodes in the Tor network. Bridge Nodes are private Tor Nodes, and are unlike other Tor Nodes, as they are not publicly listed and cannot be easily identified. Bridge Nodes essentially provide an added level of anonymity, in cases wherein publicly listed Tor Nodes are offline or actively monitored. The attacker can reverse reference the source IP address of observed TCP packets with the publicly listed Tor Node IP addresses. If the incoming IP address does not match any Node addresses, incoming connection may be from a Bridge Node, thus identifying a Bridge Node.

A malicious Sybil Node being chosen as an Entry/Guard Node or Exit Node in a Tor circuit may allow the attacker to execute Man-In-The-Middle attacks such as traffic interception and tampering.

- **Traffic Interception**

If the Exit Node is a malicious Sybil Node, after the final onion 'delaying' or decryption process with the client shared negotiated key, the attacker can potentially intercept the client's plain-text traffic during plain-text transmission with the destination, thus comprising both integrity and authenticity of the transmitted payload

4. Mitigation

The decentralized architecture of the Tor network cannot be changed in order to mitigate Sybil attacks. For any Sybil exploit mitigation strategy, the primary objective would be to limit the addition of malicious Sybil nodes to the Tor Network.

Tor Directory Authority Node's monitor Nodes, and decidedly add or remove the Nodes from the Tor consensus. The latest versions of Tor have implemented countermeasures against Sybil Nodes by enforcing IP Address restrictions allowing a single public IP Address to host only 3 Tor Nodes.

An additional mitigation technique is a change to base node selection factors of the Tor Node Selection Algorithm, which is one of the primary bases for the Sybil Attacks success, as malicious Sybil Nodes usually advertise high bandwidths in order to be selected in Tor circuit construction by the Tor Node Selection Algorithm.

3.2 Cellflood Attack

1. Description

A Cell Flood attack is essentially a Denial-Of-Service attack on the Tor network, as it essentially prevents legitimate Tor users to transmit data through the Tor Network, as the user's Tor clients will not be able to initially construct a transmission circuit through the network.

A malicious client essentially floods Tor Nodes with computationally intensive execution circuit creation requests (**CREATE**), which consequently reduces the Tor Nodes processing power, and finally results

in the Tor Nodes rejecting and dropping all circuit creation commands. As the Tor Nodes are rejecting all circuit creation commands from Tor Clients, Tor clients are essentially denied service. Additionally, the Tor network may become incapacitated, as the overloaded Tor Nodes, will be removed from the Tor consensus by the Tor Directory Authority Nodes, thus reducing the number of active online Nodes in Network. A reduction in the number of Nodes, could potentially lead to the reduction in users, which may compromise the stability and anonymity of the network.

2. Vulnerability

As discussed in previous sections, the encryption of a Tor Node in a circuit involves the Nodes public and private key. According to security researchers, the decryption process with private keys (**CREATED**) is a more computationally expensive task, as it takes 4 times longer to execute in comparison to encryption with public keys (**CREATE**).

3. Attack Vector

An attacker can initiate multiple route establishment requests by sending the Tor Nodes (**CREATE**) cell requests at a rate greater than the Tor Nodes processing rate. As the rate of receiving (**CREATE**) cells requests is larger than the Tor Nodes processing rate, the computationally intensive task will lower the bandwidth of the Node and the Node process will eventually drop (**CREATE**) cell requests by relaying (**DESTROY**) cells to the client, thus resulting in the client not being able to successfully construct a Tor transmission circuit.

4. Mitigation

A countermeasure to CellFlood attacks can be computational cryptographic puzzles. Similar to the Ethereum blockchain's proof-of-stake implementation, a Tor client will have to provide a provide computational proof to the Node in order for the Node to execute commands.

An CellFlood attack's primary objective is to essentially take the Tor Node offline using overflow attacks which will cause the Tor Node to automatically delete cells and eventually reject all incoming (**CREATE**) cell requests.

The implementation of the cryptographic puzzle essentially allows a single Tor Node to individually mitigate a CellFlood attack against itself. The puzzle countermeasure de-automates the cell deletion and demands that the malicious client will have to solve a computationally intensive cryptographic puzzle in order for the Tor Node to initiate cell deletion. The malicious client will have to allocate additional resources on the puzzle solving process, which will add a computational constraint to the (**CREATE**) cell request process, consequently disabling or slowing down the attacking malicious client. As the requests overflow has slowed, the Tor Node can process each (**CREATE**) request without significant resource constraints and additional create request reject and delete operations.

References

1. "About – Tor Metrics." Metrics.Torproject.org, metrics.torproject.org/glossary.html
2. Barbera, Marco, et al. CellFlood: Attacking Tor Onion Routers on the Cheap.
3. Bauer, Kevin, et al. Low-Resource Routing Attacks Against Tor.
4. Diego, San. USENIX Association Proceedings of the 13th USENIX Security Symposium. 2004.
5. "New Tor Denial of Service Attacks and Defenses | Tor Blog." Blog.Torproject.org, blog.torproject.org/new-tor-denial-service-attacks-and-defenses.
6. "'One Cell Is Enough to Break Tor's Anonymity' | Tor Blog." Blog.Torproject.org, blog.torproject.org/one-cell-enough-break-tors-anonymity.

7. Pries, Ryan, et al. A New Replay Attack Against Anonymous Communication Networks.
8. Winter, Philipp, et al. Identifying and Characterizing Sybils in the Tor Network Identifying and
9. “Tor: The Second-Generation Onion Router.” Svn-Archive.Torproject.org, svnarchive.torproject.org/svn/projects/design-paper/tor-design.html.